Deterministic Guardrails and Variance-Aware Evaluation for Multi-Agent Content Personalization

Submitted to *Journal of Computer Science and Technology*

Article type: Research Paper

Authors

Rishav Gupta [1,*] **Prachi Sharma** [1] **Danish Kaul** [1] **Dishant Sherawat** [1]

[1] School of Engineering and Technology, Sushant University, Gurugram, Haryana, India

[*] Corresponding author **[REPLACE: confirm who the corresponding author is]**

Contact details

Corresponding author	**[REPLACE: full name]**
Institutional e-mail	**[REPLACE: name@sushantuniversity.edu.in]**
Telephone	**[REPLACE: phone with country code]**
Postal address	**[REPLACE: full department and campus address, with PIN]**
ORCID	**[REPLACE: ORCID iD, or delete this row — free at orcid.org]**

[REPLACE: e-mail address for each remaining author, or state that correspondence should go to the corresponding author only]

Author contributions

[REPLACE: complete this honestly, e.g.: R.G. designed the agent architecture and the evaluation harness, implemented the scope and grounding guards, and ran the model comparison. P.S. implemented the frontend and the personalization panel. D.K. implemented the memory subsystem and the attachment pipeline. D.S. built the media renderer and the layout probe. All authors contributed to the manuscript and approved the submitted version.]

Declarations

Originality. This manuscript is original work. It has not been published previously and is not under consideration for publication elsewhere, in whole or in part. No version of it has been submitted to any other journal or conference.

Prior versions. None. No preliminary version of this work has appeared in any venue, so JCST's requirement to reference a preliminary version in the third person does not apply. **[REPLACE: if any part of this was presented at a college symposium, submitted as a course project report, or posted as a preprint, say so here — editors check, and an undeclared prior version is a far worse problem than a declared one]**

Competing interests. The authors declare no competing financial or non-financial interests.

Funding. **[REPLACE: state the funding source, or: This research received no specific grant from any funding agency.]**

Data and code availability. The source code, the 30-case evaluation dataset and every run report underlying the reported results will be released publicly on acceptance. They are withheld during review because the repository identifies the authors.

Ethics. This work involved no human participants, no personal data and no animal subjects. All brand and user profiles used in the evaluation are synthetic.

Use of generative AI. **[REPLACE: JCST and Springer require disclosure of generative-AI use in manuscript preparation. State plainly what was used and for what — for example that large language models were the object of study and were also used to assist drafting and code review, with all content verified by the authors. Do not omit this.]**

Anonymity statement

The manuscript file (`jcst_paper.pdf`) has been prepared for double-blind review. It contains no author name, affiliation, e-mail address, acknowledgement or repository link. Model and software names that appear in the manuscript are the subject of the study and do not identify the authors.